

Lecture 11: Priority Queue and Heap

Read: Carrano, Chpt.11.

“Good” characteristics of BST:

- Simplicity in concept & implementation.
- Objects stored in a BST can easily be sorted.
- Support general find and delete operations as well as special searchMin(Max) and deleteMin(Max) operations.
- Very efficient on average with $T_a(n) = O(\lg n)$.

“Bad” characteristics of BST:

- Worst-case complexity depends on height of tree; hence, $T_w(n) = O(n)$.
- Less efficient when many items are having identical keys.

Q: Can we design an efficient ADT with $T_w(n) = T_a(n) = O(\lg n)$?

Use Priority Queue.

ADT: Priority Queue.

A collection class in which its object has been assigned a priority with the following operations:

1. createPQ():
2. destroyPQ():
3. PQIsEmpty():
4. *PQInsert(in newItem: PQItemType)*
throw PQException
5. *PQDelete(out priorityItem: PQItemType)*
throw PQException
// delete min (or max) priority object
6. PQSize():

Remark: A PQ does not always support general delete and find operations.

Designing a Priority Queue:

Simplest Approaches:

Use a sorted or unsorted array/linked list.

Better Approach:

Use a BST.


Best Approach:

Use a heap.



Defn: A *max* (binary) *heap* H is a complete binary tree satisfying the *heap-ordered tree property* such that the priority of any node \geq priority of all its descendants.

Remark: A maximum priority element occupies the root of the max heap H.

Example: A max heap H.



Example: The following binary trees are not heaps.




Q: Can you define a *min heap*?

A *min heap* is a complete binary tree H such that the priority of any node \leq priority of all its descendants.

Implementation of Heap:

Let's consider some typical heap operation(s) to motivate our selection of data structures for the implementation of heap.

Consider inserting 9 into the following heap H.



Observe that, after insertion, we must get back a complete binary tree satisfying the heap-ordered tree property!


Q: Where should we insert the new node in order to get back a heap?

General Approach:

Always use a two-steps process:


- (1) After insertion/deletion, maintain a complete binary tree structure for H.
- (2) Re-structure H from (1) so that it will satisfy the heap-ordered tree property.

Q: Where will 9 go in order to get back a complete binary tree?




The node with key 9 must now traverse up the tree in order to get back the heap-ordered tree property.

Consider deleting max node (root) from the heap H.



Using similar two-step approach, we replace node 8 with the last node (node 5) in level-order traversal of H.




The node with key 5 must now traverse down the tree in order to get back the heap-ordered tree property.

Q: What is the most efficient data structure to implement a heap?

1. Pointer Implementation of Heap:

To facilitate the upward/downward movement in insert/delete operations, we must have parent/children information.

Consider the following node structure:



Q: Is this a good data structure?

No! This data structure is very inefficient in supporting insertion and deletion.

Why?

Remark: Do not use pointers to implement a heap!

2. Array implementation of heap:

Let H be a heap with n nodes. H can be represented using an array $A[0:\text{Max_Queue}-1]$ such that

- (1) Root of T at $A[0]$,
- (2) For any node $A[i]$, the left child (right child) of $A[i]$ is at $A[2i+1]$ ($A[2i+2]$) if exists.

Remarks:

- (1) Parent of a node $A[i]$ is at $A[(i-1)/2]$ if exists.
- (2) For $n > 1$, $A[i]$ is a leaf node in H iff $2i \geq n$.

An Alternate Array Implementation:


Skip $A[0]$ and store root of H at $A[1]$.

Hence,


- (1) Parent of node $A[i]$ is at $A[\lfloor i/2 \rfloor]$.
- (2) Left child (right child) of a node $A[i]$ is at $A[2i]$ ($A[2i+1]$).
- (3) $A[i]$ is a leaf iff $2i > n$.

Example:


H:




Array representation of H using A[0]:




Array representation of H without using A[0]:




Consider inserting 9 into a max heap H.




Inserting 9:




Compare 9 with its parent 6 and swap:



Compare 9 with its parent 7 and swap:



Compare 9 with its parent 8 and swap:




Since node 9 has no parent; process terminates.

In general, for inserting a new item X into H , we need to find a location to insert X along the path from X (after insertion) to the root of H by *repeatedly comparing X with his parent, grandparent, ..., until either a node with priority $\geq X$ is found or the root is reached.*


Remark: Insert (remove) operation can be somewhat speeded up if we will attempt to find the *final* location for X first instead of placing X into various locations in H only find out later that it needs to be moved to another location again.

Example: Consider inserting 9 into the heap H.


Creating new location for 9:




Compare 9 with its parent 6; 6 moves down:




Compare 9 with its parent 7; 7 moves down:



Compare 9 with its parent 8; 8 moves down:



Insert 9 into its final location; process terminates:



Using Array implementation:

Using Array implementation:

0	1	2	3	4	5	6	7	8	9	10	...		
8	7	5	5	6	3	2	1	3	5				


0	1	2	3	4	5	6	7	8	9	10	...		
8	7	5	5		3	2	1	3	5	6			

0	1	2	3	4	5	6	7	8	9	10	...		
8		5	5	7	3	2	1	3	5	6			

0	1	2	3	4	5	6	7	8	9	10	...		
	8	5	5	7	3	2	1	3	5	6			


0	1	2	3	4	5	6	7	8	9	10	...		
9	8	5	5	7	3	2	1	3	5	6			

Consider deleting the highest priority item (root) from the original heap H.




Q: What happens after 8 is removed?


Replace the root of H with the last item (in level order) 5.




Compare 5 with its two children, swap with 7:




Compare 5 with its two children, swap with 6:



Insert 5 into its final location; process terminates:





In general, we replace the root (highest priority item) of H with the “last” item X in H . We then *repeatedly compare X with its child (children), swap with the larger child if necessary, until $X \geq$ its child (children if X has two children) or a leaf is reached*.

Performance analysis:

For insert/remove operations, in the worst-case, we have to either traverse all the way from a leaf to the root (for insert) or traverse all the way from the root to a leaf (for delete). Hence, the complexities of both algorithms will depend on the height of the heap.

Since a heap is a complete binary tree and a complete binary tree with n nodes has height $\lfloor \lg n \rfloor + 1$, $T_w(n) = O(\lg n)$, where n is the number of items in the heap.

Array implementation:

0	1	2	3	4	5	6	7	8	9	...
8	7	5	5	6	3	2	1	3	5	

0	1	2	3	4	5	6	7	8	9	...
	7	5	5	6	3	2	1	3		

0	1	2	3	4	5	6	7	8	9	...
7		5	5	6	3	2	1	3		

0	1	2	3	4	5	6	7	8	9	...
7	6	5	5		3	2	1	3		

0	1	2	3	4	5	6	7	8	9	...
7	6	5	5	5	3	2	1	3		


Q: How do we build the initial heap H?
Two build-heap methods:

1. **Top-down approach:** Insert items in the set, one by one, into an initially empty heap.

$$T_w(n) = O(n \lg n).$$

2. **Bottom-up approach:** First form a complete binary tree H for S according to its given order. If we scan the nodes of H in the reverse level order (leaf-to-root), observe that a leaf by itself is a heap and two heaps can be combined together with their common parent to form a bigger heap (as in delete operation), we can grow a heap for S in a bottom-up fashion.


Heapify as in delete operation by moving a root node down the tree so as to satisfy the heap-order tree property:



$$T_w(n) = O(n).$$


Example: Build a heap for $S = \{5, 3, 2, 6, 8, 5, 7, 1, 3, 5\}$.

Form initial complete binary tree:




Since $n = 10$, first non-leaf node needs to be checked has array index $\lfloor n/2 \rfloor - 1 = 4$, follows by nodes with index 3, 2, 1, 0.


For $A[2]$, compare 2 with 5 and 7, swap with 7:




For A[1], compare 3 with 6 and 8, swap with 8:




Compare 3 with 5, swap with 5:



Compare 5 with 8 and 7, swap with 8:



Compare 5 with 6 and 5, swap with 6:



Compare 5 with 1 and 3, terminates process!

HW: Repeat above constructions using array.

PQ Sorting Revisited:

1. Build a PQ Q for S.
2. Delete repeatedly until Q is empty.

Heap Sort:

1. Build a max-heap H for S.
2. Deletemax repeatedly until H is empty. (Swap max item with the current last item in array.)

Worst-Case Complexity of Heap Sort:

Building a heap: $\Theta(n)$.

DeleteMax: $\Theta(\lg n)$.

Hence, $T(n) = \Theta(n \lg n)$

Worst-Case Time Comparison of PQ Implementations:

Assuming max-heap:

<u>PQ Operation</u>	<u>Heap</u>	<u>BST</u>	<u>Sorted List</u>	<u>Unsorted Array</u>
<i>Build/Organize</i>	$O(n)$	$O(n^2)$	$O(n^2)$	$O(n)$
<i>Insert</i>	$O(\lg n)$	$O(n)$	$O(n)$	$O(1)$
<i>Search</i>	$O(n)$	$O(n)$	$O(n)$	$O(n)$
<i>GeneralDelete</i>	$O(n)$	$O(n)$	$O(n)$	$O(n)$
<i>DeleteMax</i>	$O(\lg n)$	$O(n)$	$O(n)$	$O(n)$
<i>DeleteMin</i>	$O(n)$	$O(n)$	$O(1)$	$O(n)$

Extension: k-heap, $k \geq 3$.

2-heap

Complete binary tree

Heap-ordered tree

k-heap

Complete k-ary tree

Heap-ordered tree

Implementation:

Array implementation with root at $A[1]$.

Parent of $A[i]$ at $A[(i+k-2)/k]$,

1st child at $A[ki-k+2]$,

2nd child at $A[ki-k+3]$,

3rd child at $A[ki-k+4]$,

...

j^{th} child at $A[ki-k+j+1]$, $1 \leq j \leq k$, if exists.

Example:

Given $A[i]$ in a 3-heap.

Parent of $A[i]$ at $A[(i+1)/3]$,

1st child at $A[3i-1]$,

2nd child at $A[3i]$,

3rd child at $A[3i+1]$, if exists.

HW: Compute parent-children info if rooted at $A[0]$.